

The faculty your agent is missing

LLMs gave your agent reasoning. RAG gave it memory.
Aito gives it **intuition**.

A live agent that calls Aito ops as tools — grounded numbers it can't invent, and the action that wins. Three working surfaces (sales, company 360, support), one predictive database, no model training. Better, faster, cheaper — and higher-yield.

**Predictive Agent**
Powered by Aito.ai

START HERE

Overview

NORTHLIGHT · SALES

Sales agent

Opportunity Assistant

Toolbox

NORTHWIND CLOUD · COMPANY

Company AI agent

Toolbox

SONIPRA TELECOM · SUPPORT

Ticket resolution

Tool routing · short-list

Human handoff

analyze · assist · automate
every view is live Aito + gpt-5-mini

THE FACULTY YOUR AGENT IS MISSING

LLMs gave your agent reasoning. RAG gave it memory. Aito gives it intuition.

A neural network turns experience into instant answers — but that intuition is **frozen at training time**, and it's about the whole internet, not your business. LLMs are brilliant **amnesiacs**: they don't know your customers, and training one on your data isn't feasible. Aito does the same thing — pattern into answer — **live, over your own data, with no training**.

Ask it about your customers, orders, tickets, codes — and it just **knows**, with a calibrated sense of how sure it is. The **known and the unknown**, through one door.

Meet the agent

How it fits

same act as the neural net

but live · no training · no MLOps

over your data, not the world's

calibrated \$p + \$why

WHAT IT IS

Reasoning · Memory · Intuition

An agent needs all three. You already have two. Aito is the third — the **same kind of pattern-machine as the model**, specialized to your data: it turns what you've seen into an instant, calibrated answer, with no training and nothing to forget.

aito..

THE PREDICTIVE DB

5

\$p

\$why

CORE OPS

CALIBRATED

EXPLAINED

one query · no training

Aito turns your rows into an **instant, calibrated answer** — the same act as a neural net, but live and over **your** data, with nothing to train. Every demo in the menu is a real call to one of five ops.

THE FIVE OPS

```
_predict  class + $p + $why
_match    relevant memory
_relate   drivers / lift
_estimate a number
_recommend best next action
```

LEARN MORE

[View live schema \(JSON\)](#)

[Predict API reference](#)

[Source on GitHub](#)

Start free trial →

Where a capable agent quietly falls down

None of these mean you picked the wrong model or built the wrong platform. They're the predictable failure modes of asking one LLM to reason *and* remember *and* do arithmetic over a large, structured, ever-changing dataset. Each has a one-query fix.



Tool / option sprawl

Hundreds of tools or SKUs in context → selection degrades, prompts bloat.



An LLM call on every step

Multi-step workflows take seconds and burn tokens, per ticket, at scale.



Vector search misfires

Embeddings dilute identifiers — the nearest neighbour is the wrong customer.



Bad with numbers

Aggregation, drivers, estimates — the model guesses, often confidently wrong.



No sense of "how sure"

Overconfident output gives no signal for when to act vs ask a human.



Amnesiac about your data

It knows the internet, not your customers — and you can't feasibly train it on them.

The solution: Aito is a tool in the agent's toolbox

Aito is a predictive database. Load your history; query for predictions, recommendations and statistics through SQL-like calls. **No model training, no retraining schedule, no MLOps.** Every answer carries a calibrated probability (\$p) and an explanation (\$why) that traces straight to your data.

The agent keeps doing the reasoning and the talking. When it needs a number it can't invent — win odds, effort, churn, the action that wins — it **calls an Aito op as a function**. The same agent with the Aito tools switched off has to guess, and says so. That is the whole pitch, A/B in one toggle.

_predict — a class + calibrated \$p + \$why
_match — the relevant memory, by structure
_relate — drivers / lift behind a number
_estimate — a number, from history
_recommend — the next action that maximises a goal

One query, any language. Call it like a tool or an MCP endpoint from any agent framework. Your platform stays the brain; Aito is the instant, calibrated memory underneath it.

Zero MLOps. No model files, no retrain, no drift. A row added today is in the next prediction — which is also how the agent *learns*: log an outcome and the very next recommendation is sharper.

Calibrated & explainable. Every answer has a \$p and a \$why. Auditable by design — the \$p is a governance gate: auto when sure, escalate when not.

Aito in the toolbox — the sales agent

A live gpt-5-mini chat loop. Ask it about a deal; it reasons, then calls Aito ops for the numbers it can't invent: win-odds (`_predict + $why`), effort (`_estimate`), reference projects (`_query`), and the outreach that books the most meetings (`_recommend`) — quoting the **lift over the unoptimised baseline**. Every tool call is rendered inline, so you watch Aito work inside the agent. Money/state actions are gated: it drafts the email, it never sends.

The screenshot displays the Predictive Agent interface. On the left is a dark sidebar with navigation links: Overview, NORTHLIGHT · SALES, Sales agent (highlighted), Opportunity Assistant, Toolbox, NORTHWIND CLOUD · COMPANY, Company AI agent, 360 Dashboard, Toolbox, SONIPRA TELECOM · SUPPORT, Ticket resolution, Tool routing · short-list, and Human handoff. The main area shows a chat window with a message: "Warm referral into an Enterprise SaaS company that wants a data platform build, sole-source. Should we pursue it, and how big is the job?". Below this, a "CALLED THE TOOLBOX" section lists several tool calls: `_predict win_odds` (181ms) resulting in "win 90% · 3 drivers"; `_estimate estimate_effort` (183ms) resulting in "84 person-days"; `_query find_references` (56ms) resulting in "3 won references"; and `_recommend recommend_outreach` (265ms) resulting in "Warm intro · Case study · meeting 59% · 3x vs 20% baseline". An `action propose_send_email` (gated) is shown with a draft email body. At the bottom is an input field "Ask a question..." and a "Send" button. On the right, a dark panel titled "aito.. AITO IN THE TOOLBOX" shows three modes: assist (OPTIMIZE), +yield (OUTCOME), and live (GPT-5-MINI). It includes a description of the tool's purpose, a code snippet for `recommend_outreach`, and links to learn more.

Toolbox on/off · grounded numbers vs flagged guesses · 4 Aito calls in 0.8 s, \$0 · the LLM call reasons; Aito supplies the facts.

A 360° copilot that optimises KPIs — the company agent

One linked customers master joined to sales, support, product, finance and CX. A SQL+LLM BI bot can COUNT (*); this copilot calls Aito for the **360 KPI snapshot** (conversion, churn, NPS, CSAT, adoption, on-time revenue), the **single lever that moves each one** with the projected lift, and a **customer-360 join** across every domain. It doesn't just report — it optimises, drafts the play (gated), and learns: Aito has no training step, so a logged outcome sharpens the next prediction.



Predictive Agent

Powered by Aito.ai

START HERE

Overview

NORTHLIGHT · SALES

Sales agent

Opportunity Assistant

Toolbox

NORTHWIND CLOUD · COMPANY

Company AI agent

360 Dashboard

Toolbox

SONIPRA TELECOM · SUPPORT

Ticket resolution

Tool routing · short-list

Human handoff

analyze · assist · automate

every view is live Aito + gpt-5-mini

Aito tools ON

6/6 tools · set in Toolbox

2 Aito calls · 7571 tokens · \$0.0030

Give me the 360 KPI snapshot for our SMB Free-plan customers, then the single biggest lever to cut their churn and the projected impact.

CALLLED THE TOOLBOX

_predict

kpi_snapshot

439ms

size=SMB · plan=Free

→ conv 41% · churn 35% · NPS det 45% · CSAT 56% · adopt 49%

_predict · _relate · _recommend

optimize_kpi

577ms

kpi=churn · size=SMB · plan=Free

→ Churn: 35% → 23% (12pp) via CSM motion=Exec-sponsor

action

launch_play

gated

KPI churn

SEGMENT SMB — Free plan

PLAY

Move segment from no CSM to an Exec-sponsor CSM motion: assign Exec-level sponsored CSMs to high-risk Free-plan SMB accounts (health=Red or onboarding=None). Prioritise accounts with recent product activity then no onboarding. Outreach plan: 1) 1:1 kickoff within 7 days, 2) weekly 15-min adoption check-ins for 6 weeks, 3) playbook to convert to Starter on value signals. Track churn outcome at 90 days.

EXPECTED_IMPACT

Churn 35% → 23% (-12 percentage points) projected by model; lever lift multiplier ~1.31.

Draft only — never acts automatically

Approve play

Snapshot — SMB Free-plan

- Conversion: 41%

- Churn: 35%

- NPS: detractor band

- CSAT: good

Ask a question...

Send

aito..

360 · OPTIMISE KPIS

360

optimise

learn

LINKED

ANY KPI

NO RETRAIN

see · optimise · act · learn

One **linked** customers master across sales, support, product, finance and CX. A SQL+LLM bot can COUNT(*); this agent calls Aito for the **360 KPIS**, the **lever that moves each one** (_recommend + projected lift), and drafts the play. Aito has no training step — a logged outcome sharpens the next prediction. A closed loop, no retrain.

OPTIMISE A KPI

```
optimize_kpi("churn",
  {size:"SMB", plan:"Free"})
// = _predict $why + _recommend
// 65% → 78% via Exec-sponsor CSM
```

LEARN MORE

View live schema (JSON)

Predict API reference

Source on GitHub

Start free trial →

“Churn 40% → 25% via an Exec-sponsor CSM motion” — driver, lever and projected lift, computed live from the company’s own linked data. See · optimise · act · learn.

Aito Predictive Agent · agent.aito.ai · Apache 2.0

4

The 360 Dashboard — every KPI, diagnosed and prescribed

The same ops, called **directly** (no agent) as a data view. Pick a customer segment and each KPI shows three things, each with a “?” that opens the live explanation — the per-prediction \$why pattern used across the Aito demos:

1234

Why the rate is the rate

_predict \$why: the base rate scaled by each segment attribute's lift — e.g. conversion base 40% × SMB × 1.11 × Free × 0.92.




Root causes

_relate scoped to the segment (\$on): within SMB-Free, Red-health customers churn 44% vs 28% → ×1.39. Drivers (>1) and protective factors (<1).



Levers + projected lift

_recommend ranks the actionable lever; the recommendation differs by segment (SMB → Guided trial, Enterprise → CSM-led) and projects the result.



Diagnosis (_relate) and prescription (_recommend) on every KPI, each explained by a live \$why — the predictive database as a closed optimise-act-learn loop.

Predict-first support — the resolution console

The same ticket, two engines, side by side. An LLM agent reasons the resolution out on every ticket — seconds and tokens. `aito._predict` reads the intent and the one parameter it needs straight from history — two calls, sub-second, \$0, with a calibrated \$why. A confident hit is served instantly (a cache hit); a miss falls through to the LLM, whose answer becomes the next cache entry. The \$p gate routes the unsure ones to a human, and anything sensitive (refund, cancel) is gated regardless.

The screenshot displays the Predictive Agent interface, comparing the performance of AITO and LLM Agent in resolving a ticket. The interface is divided into several sections:

- Header:** Predictive Agent Powered by Aito.ai. Navigation links include Overview, Sales agent, Opportunity Assistant, and Toolbox.
- Performance Metrics:**
 - AITO - PREDICT-FIRST:** 229ms, two _predict calls, \$0 LLM.
 - LLM AGENT - LIVE:** 2.2s, gpt-5-mini, one call.
 - SPEED-UP:** Aito vs one agent call.
 - LLM COST / RESOLUTION:** Aito \$0 per ticket.
- Same ticket, two engines:** The LLM agent reasons the resolution out on every ticket — seconds and tokens. `aito._predict` reads the **intent** and the one parameter it needs straight from history — two calls, sub-second, \$0, with a calibrated **why**. Watch the response rates live.
- Ticket Details:** Outage - Helsinki. Repair - cracked screen. Cancel - broadband. Refund - roaming. Balance.
- Resolution Query:** Is there a network outage? Nothing works in Helsinki. alerts.monitoring.io
- LLM Agent:** 2.2s. Reasoning: gpt-5-mini reasoning...
- Predictive Layer:** Aito_predict 229ms. Return outage status for Helsinki. Intent: check_outage (0.98). Location: Helsinki (0.98). Why - verifiable from history. Base Probability: Historical rate for check_outage (17%). Pattern Match: When text is Is there a network outage? Nothing works in Helsinki. (5.7x). 17% x 5.7 = 98%. 0 model calls, \$0, calibrated \$p.
- aito.. THE PREDICTIVE DB - PREDICT:** 229ms AITO LATENCY. intent PREDICT FIELD. live LLM LATENCY. `_predict`. Aito is a **predictive cache in front of the LLM**. Each ticket is resolved by reading the **intent** and the one parameter it needs from history — **two _predict calls, no chain**. A confident hit fires instantly and free; a miss falls through to the LLM, whose answer becomes the next cache entry. LIVE QUERY: POST /api/v1/_predict { "from": "resolutions", "where": { "text": "_ticket_", "sender_domain": "_" }, "predict": "intent", "select": ["\$p", "\$why"] } // - check outage (p = 0.98). VERIFY YOURSELF: Every routed call traces back to an Aito query — no model file, no retrain. A row added today is in the next prediction. LEARN MORE: View live schema (JSON), Predict API reference, Source on GitHub. Start free trial.

Benchmarked, not asserted

Three failure modes every agent team hits — each run as a real benchmark (live Aito

1. live gpt-5-mini on seeded, realistic data) against the tool a good engineer would

otherwise reach for.

01 Shortlisting holds at scale

Embedding-retrieval shortlist degrades as the catalog grows; Aito holds and hands the LLM ~16× fewer tokens for the same pick.

02 ~9–10× lower latency

A 6-step LLM resolution chains to ≈22 s; Aito predicts in parallel in ~0.15 s — resolved before the agent clears step one.

03 Right context-memory

Vector search picks the wrong customer's memory 86% of the time; Aito conditions on structure and recovers what the text can't identify.

See it just know — all three agents, live.

Real predictions, real latency, real cost — at agent.aito.ai. A primitive your agents call, not another platform to adopt.